

## Lesson-03 :: Group Functions, Joins, Sub-Queries

---

### Group Functions

---

AKA -           Aggregate Functions  
          Multi-row Functions

They work on a group of rows.  
If the grouping is not explicitly specified, all the  
rows of the table are taken into account.

### Examples

```
SQL> SELECT SUM( salary ) FROM employees;
```

```
SUM(SALARY)
-----
      691416
```

```
SQL> SELECT AVG( salary ) FROM employees;
```

```
AVG(SALARY)
-----
    6461.83178
```

```
SQL> SELECT MIN( salary ) FROM employees;
```

```
MIN(SALARY)
-----
       2100
```

```
SQL> SELECT MAX( salary ) FROM employees;
```

```
MAX(SALARY)
-----
      24000
```

```
SQL> SELECT COUNT( commission_pct ) FROM employees;
```

```
COUNT(COMMISSION_PCT)
-----
                36
```

NOTE: With COUNT() if a column name is specified, it returns  
the count of Non-NULL values.

```
SQL> SELECT COUNT( * ) FROM employees;
```

```
COUNT(*)
-----
      107
```

### Using GROUP BY Clause

---

The GROUP BY clause enables grouping of rows.

```
SQL> SELECT department_id, SUM( salary ) FROM employees
2  GROUP BY department_id
3  ORDER BY department_id;
```

DEPARTMENT_ID	SUM(SALARY)
10	4400
20	19000
30	24900
40	6500
50	156400
60	28800
70	10000
80	304500
90	58000
100	51608
110	20308

```
SQL> SELECT department_id, SUM( salary ), AVG( salary ), MIN( salary ),
MAX( salary ) FROM employees
2  GROUP BY department_id
3  ORDER BY department_id;
```

DEPARTMENT_ID	SUM(SALARY)	AVG(SALARY)	MIN(SALARY)	MAX(SALARY)
10	4400	4400	4400	4400
20	19000	9500	6000	13000
30	24900	4150	2500	11000
40	6500	6500	6500	6500
50	156400	3475.55556	2100	8200
60	28800	5760	4200	9000
70	10000	10000	10000	10000
80	304500	8955.88235	6100	14000
90	58000	19333.3333	17000	24000
100	51608	8601.33333	6900	12008
110	20308	10154	8300	12008
	7000	7000	7000	7000

#### The HAVING Clause

=====

Used to filter group(s) based on a condition.

```
1  SELECT department_id, SUM( salary ), AVG( salary ), MIN( salary ),
MAX( salary ) FROM employees
2  GROUP BY department_id
3  HAVING AVG(salary) >= 10000
4* ORDER BY department_id
SQL> /
```

DEPARTMENT_ID	SUM(SALARY)	AVG(SALARY)	MIN(SALARY)	MAX(SALARY)
70	10000	10000	10000	10000
90	58000	19333.3333	17000	24000
110	20308	10154	8300	12008

Example:

```
SQL> -- Getting a HEAD COUNT department-wise
```

```
SQL> ed
```

Wrote file afiedt.buf

```
1  SELECT department_id, COUNT( * ) "No of Employees" FROM employees
2  GROUP BY department_id
3* ORDER BY 1
SQL> /
```

DEPARTMENT_ID	No of Employees
10	1
20	2
30	6
40	1
50	45
60	5
70	1
80	34
90	3
100	6
110	2

Get the group which has employees more than or equal to 30

```
SQL> ed
```

Wrote file afiedt.buf

```
1  SELECT department_id, COUNT( * ) "No of Employees" FROM employees
2  GROUP BY department_id
3  HAVING COUNT( * ) >= 30
4* ORDER BY 1
SQL> /
```

DEPARTMENT_ID	No of Employees
50	45
80	34

Getting a count of employees based on the job id. Moreover the count should be greater than or equal to FIVE.

```
1  SELECT job_id, COUNT( * ) FROM employees
2  GROUP BY job_id
3* HAVING COUNT( * ) >= 5
SQL> /
```

JOB_ID	COUNT(*)
FI_ACCOUNT	5
IT_PROG	5
PU_CLERK	5
SA_MAN	5
SA_REP	30
SH_CLERK	20
ST_CLERK	20
ST_MAN	5

Example:

Getting a count of 'first\_name' which occurs more than once

```
1 SELECT first_name, COUNT( * ) FROM employees
2 GROUP BY first_name
3* HAVING COUNT(*) > 1
SQL> /
```

FIRST_NAME	COUNT(*)
Peter	3
Michael	2
Steven	2
John	3
Julia	2
William	2
Karen	2
Kevin	2
David	3
Jennifer	2
Randall	2

Example:

Getting a count of 'first\_name' more than once, but only from department 50 or 80

```
1 SELECT first_name, COUNT( * ) FROM employees
2 WHERE department_id IN (50, 80)
3 GROUP BY first_name
4* HAVING COUNT(*) > 1
SQL> /
```

FIRST_NAME	COUNT(*)
Peter	3
Julia	2
John	2
Kevin	2
Randall	2
David	2
James	2

#####

Why split data in Multiple-Tables?

~~~~~

RDBMS Concepts

- Normalization
- Minimizing Data Redundancy
- Data Consistency

We deliberately split the data in multiple tables.

What is Joining?

=====

Joining is ability of the SELECT statement to access the data existing in two or more tables.

Oracle supports

- Oracle proprietary joins
- SQL 1999 compliant joins

Cartesian Product

-----  
It combines the number of rows of first table with the number of rows of the second table and gives a complete list.

i.e. If Table-A has 5 rows and Table-B has 3 row, the Cartesian Product will generate 15 rows.

This will happen when we do not have a VALID where clause.

Example-1: Oracle proprietary

```
SQL> SELECT department_id, city
      2 FROM departments, locations;
```

NOTE: The WHERE clause is NOT SPECIFIED in this case.

Example-2: Using SQL 1999 compliant

```
SELECT department_id, city
FROM departments
CROSS JOIN locations;
```

Types of Joins

-----

[1] Equi-Join

We make use of the EQUALITY Operator (i.e. = )

[2] Outer Join

Right Outer Join

Left Outer Join

Full Outer Join

[3] Non-Equi Join

We make use of Operator other than the Equality (i.e. = )

[4] Self Join

If a table joins to itself it is called Self Join

It is also called as Recursive Relationship.

EquiJoin

=====

Example-1a

```
SQL> SELECT department_name, city
      2 FROM departments, locations
      3 WHERE departments.location_id = locations.location_id;
```

| DEPARTMENT_NAME | CITY                |
|-----------------|---------------------|
| IT              | Southlake           |
| Shipping        | South San Francisco |
| Administration  | Seattle             |
| Purchasing      | Seattle             |

```

:
:
Payroll                Seattle
Marketing              Toronto
Human Resources        London
Sales                  Oxford
Public Relations       Munich

```

Example-1b : SQL 1999 Compliant

SQL> ed

Wrote file afiedt.buf

```

1  SELECT department_name, city
2  FROM departments
3* NATURAL JOIN locations

```

Example-2a - Getting the employee first name and department name

SQL> SELECT first\_name, department\_name

2 FROM employees, departments

3 WHERE employees.department\_id = departments.department\_id;

| FIRST_NAME | DEPARTMENT_NAME |
|------------|-----------------|
| Jennifer   | Administration  |
| Pat        | Marketing       |
| Michael    | Marketing       |
| Sigal      | Purchasing      |

Example-2b - Getting the employee first name and department name

Example-3a : Getting employee first name and job title

SQL> SELECT first\_name, job\_title

2 FROM employees, jobs

3 WHERE employees.job\_id = jobs.job\_id;

| FIRST_NAME  | JOB_TITLE                     |
|-------------|-------------------------------|
| William     | Public Accountant             |
| Shelley     | Accounting Manager            |
| Jennifer    | Administration Assistant      |
| Steven      | President                     |
| Neena       | Administration Vice President |
| Lex         | Administration Vice President |
| Jose Manuel | Accountant                    |
| Ismael      | Accountant                    |
| :           |                               |
| :           |                               |

Example-3b : Getting employee first name and job title

SELECT first\_name, job\_title

FROM employees

NATURAL JOIN jobs

## Non-Equi Join

=====

### Example-1

```
SQL> SELECT ename, sal, grade
  2   FROM emp, salgrade
  3  WHERE sal >= losal AND sal <= hisal;
```

| ENAME  | SAL  | GRADE |
|--------|------|-------|
| SMITH  | 800  | 1     |
| JAMES  | 950  | 1     |
| ADAMS  | 1100 | 1     |
| WARD   | 1250 | 2     |
| MARTIN | 1250 | 2     |
| MILLER | 1300 | 2     |
| TURNER | 1500 | 3     |
| ALLEN  | 1600 | 3     |
| CLARK  | 2450 | 4     |
| SCOTT  | 3000 | 4     |
| FORD   | 3000 | 4     |
| KING   | 5000 | 5     |

```
  1  SELECT ename, sal, grade
  2  FROM emp, salgrade
  3* WHERE sal BETWEEN losal AND hisal;
```

## Outer Join

=====

### Example

```
SQL> SELECT ename, dname FROM emp, dept
  2  WHERE emp.deptno = dept.deptno;
```

| ENAME  | DNAME      |
|--------|------------|
| CLARK  | ACCOUNTING |
| KING   | ACCOUNTING |
| MILLER | ACCOUNTING |
| JONES  | RESEARCH   |
| FORD   | RESEARCH   |
| ADAMS  | RESEARCH   |
| SMITH  | RESEARCH   |
| SCOTT  | RESEARCH   |
| WARD   | SALES      |
| TURNER | SALES      |
| ALLEN  | SALES      |
| JAMES  | SALES      |
| BLAKE  | SALES      |
| MARTIN | SALES      |

14 rows selected.

However, there is a department by name OPERATIONS for which there are no employees.

To get the department name for which there are no employees we use the OUTER JOIN concept.

#### Example-1a

```
1  SELECT ename, dname FROM emp, dept
2*  WHERE emp.deptno(+) = dept.deptno;
SQL> /
```

| ENAME  | DNAME      |
|--------|------------|
| CLARK  | ACCOUNTING |
| KING   | ACCOUNTING |
| MILLER | ACCOUNTING |
| JONES  | RESEARCH   |
| FORD   | RESEARCH   |
| ADAMS  | RESEARCH   |
| SMITH  | RESEARCH   |
| SCOTT  | RESEARCH   |
| WARD   | SALES      |
| TURNER | SALES      |
| ALLEN  | SALES      |
| JAMES  | SALES      |
| BLAKE  | SALES      |
| MARTIN | SALES      |
|        | OPERATIONS |

NOTE: The (+) syntax for getting Outer Join is Oracle property.

#### Example-1b : Performing RIGHT OUTER JOIN

```
SELECT ename, dname
FROM emp e
RIGHT OUTER JOIN dept d
ON (e.deptno = d.deptno)
```

#### Self Join

=====

If a table joins to ITSELF then we say we are performing SELF JOIN

```
SQL> -- The employee name and the manager to whom they report
```

```
SQL>
```

```
SQL> SELECT worker.ename "Employee Name", manager.ename "Manager"
2  FROM emp worker, emp manager
3  WHERE worker.mgr = manager.empno;
```

#### Employee N Manager

| Employee | Manager |
|----------|---------|
| FORD     | JONES   |
| SCOTT    | JONES   |
| TURNER   | BLAKE   |
| ALLEN    | BLAKE   |
| WARD     | BLAKE   |
| JAMES    | BLAKE   |
| MARTIN   | BLAKE   |
| MILLER   | CLARK   |
| ADAMS    | SCOTT   |
| BLAKE    | KING    |
| JONES    | KING    |
| CLARK    | KING    |



Performing an LEFT Outer Join to know which employee does not have a Manager

```
1 SELECT worker.ename "Employee Name", manager.ename "Manager"
2 FROM emp worker, emp manager
3* WHERE worker.mgr = manager.empno(+)
SQL> /
```

Employee N Manager

| Employee | Manager |
|----------|---------|
| FORD     | JONES   |
| SCOTT    | JONES   |
| JAMES    | BLAKE   |
| TURNER   | BLAKE   |
| MARTIN   | BLAKE   |
| WARD     | BLAKE   |
| ALLEN    | BLAKE   |
| MILLER   | CLARK   |
| ADAMS    | SCOTT   |
| CLARK    | KING    |
| BLAKE    | KING    |
| JONES    | KING    |
| SMITH    | FORD    |
| KING     |         |

14 rows selected.

Using SQL 1999 complaint

```
SELECT worker.ename "Employee", manager.ename "Manager"
FROM emp worker
LEFT OUTER JOIN emp manager
ON ( worker.mgr = manager.empno)
/
```

We can even perform the same with 'HR' schema.

```
SQL> SELECT w.first_name "Employee", m.first_name "Manager"
2 FROM employees w, employees m
3 WHERE w.manager_id = m.employee_id;
```

| Employee  | Manager |
|-----------|---------|
| William   | Gerald  |
| Lisa      | Gerald  |
| Sundita   | Gerald  |
| Tayler    | Gerald  |
| Harrison  | Gerald  |
| Elizabeth | Gerald  |
| Alexander | Lex     |
| Clara     | Alberto |
| Mattea    | Alberto |
| David     | Alberto |
| Danielle  | Alberto |
| Amit      | Alberto |
| Sundar    | Alberto |
| Nandita   | Adam    |

```
TJ                Adam
:
:
```

```
1  SELECT w.first_name "Employee", m.first_name "Manager"
2  FROM employees w, employees m
3* WHERE w.manager_id = m.employee_id(+)
SQL> /
```

| Employee  | Manager |
|-----------|---------|
| Sundita   | Gerald  |
| Elizabeth | Gerald  |
| William   | Gerald  |
| Tayler    | Gerald  |
| Harrison  | Gerald  |
| Lisa      | Gerald  |
| Alexander | Lex     |
| Amit      | Alberto |
| :         |         |
| :         |         |
| Alyssa    | Eleni   |
| Ellen     | Eleni   |
| Steven    |         |

NOTE: Steven does not have a Manager.

SQL 1999 - Join Syntax  
=====

[1] Cross Join

Creates a Cartesian Product

Example

```
SQL> ed
Wrote file afiedt.buf
```

```
1  SELECT department_id, city
2  FROM departments
3* CROSS JOIN locations
SQL> /
```

[2] Natural Join

Is the Equi Join

```
SQL> SELECT department_name, city
2  FROM departments
3  NATURAL JOIN locations;
```

| DEPARTMENT_NAME | CITY                |
|-----------------|---------------------|
| IT              | Southlake           |
| Shipping        | South San Francisco |
| Administration  | Seattle             |
| Purchasing      | Seattle             |
| :               |                     |

```

:
Payroll                Seattle
Marketing              Toronto
Human Resources        London
Sales                  Oxford
Public Relations       Munich

```

#### [2-A] Natural Join - Using the USING Clause

```

SQL> ed
Wrote file afiedt.buf

  1  SELECT department_name, city
  2  FROM departments
  3  JOIN locations
  4* USING(location_id, location_id)
SQL> /

```

NOTE: The USING clause should be preferred with the column names do not match along with their respective data types.

#### [2-B] Natural Join - Using the ON Clause

```

SQL> ed
Wrote file afiedt.buf

  1  SELECT department_name, city
  2  FROM departments
  3  JOIN locations
  4* ON (departments.location_id = locations.location_id)
SQL> /

```

```

SQL> ed
Wrote file afiedt.buf

  1  SELECT department_name, city
  2  FROM departments
  3  JOIN locations
  4  ON (departments.location_id = locations.location_id)
  5* AND city IN ('Southlake', 'Oxford', 'Toronto', 'Munich')
SQL> /

```

| DEPARTMENT_NAME  | CITY      |
|------------------|-----------|
| Public Relations | Munich    |
| Sales            | Oxford    |
| IT               | Southlake |
| Marketing        | Toronto   |

#### [3] Outer Join

```

SQL> ed
Wrote file afiedt.buf

  1  SELECT ename, dname
  2  FROM emp RIGHT OUTER JOIN dept
  3* ON (emp.deptno = dept.deptno )
SQL> /

```

| ENAME  | DNAME      |
|--------|------------|
| CLARK  | ACCOUNTING |
| KING   | ACCOUNTING |
| MILLER | ACCOUNTING |
| JONES  | RESEARCH   |
| FORD   | RESEARCH   |
| ADAMS  | RESEARCH   |
| SMITH  | RESEARCH   |
| SCOTT  | RESEARCH   |
| WARD   | SALES      |
| TURNER | SALES      |
| ALLEN  | SALES      |
| JAMES  | SALES      |
| BLAKE  | SALES      |
| MARTIN | SALES      |
|        | OPERATIONS |

Example

```
SQL> SELECT ename, dname
      2 FROM new_emp LEFT OUTER JOIN dept
      3 ON (new_emp.deptno = dept.deptno);
```

| ENAME  | DNAME      |
|--------|------------|
| MILLER | ACCOUNTING |
| KING   | ACCOUNTING |
| CLARK  | ACCOUNTING |
| FORD   | RESEARCH   |
| ADAMS  | RESEARCH   |
| SCOTT  | RESEARCH   |
| JONES  | RESEARCH   |
| SMITH  | RESEARCH   |
| JAMES  | SALES      |
| TURNER | SALES      |
| BLAKE  | SALES      |
| MARTIN | SALES      |
| WARD   | SALES      |
| ALLEN  | SALES      |

Harshan

Example : Full Outer Join

```
SQL> ed
Wrote file afiedt.buf

      1 SELECT ename, dname
      2 FROM new_emp FULL OUTER JOIN dept
      3* ON (new_emp.deptno = dept.deptno)
SQL> /
```

| ENAME | DNAME    |
|-------|----------|
| SMITH | RESEARCH |
| ALLEN | SALES    |
| WARD  | SALES    |
| JONES | RESEARCH |

|         |            |
|---------|------------|
| MARTIN  | SALES      |
| BLAKE   | SALES      |
| CLARK   | ACCOUNTING |
| SCOTT   | RESEARCH   |
| KING    | ACCOUNTING |
| TURNER  | SALES      |
| ADAMS   | RESEARCH   |
| JAMES   | SALES      |
| FORD    | RESEARCH   |
| MILLER  | ACCOUNTING |
| Harshan |            |
|         | OPERATIONS |

Observe the output, we get those employee(s) who are not mapped to a department, plus we get those department(s) for which we do not have employees.

NOTE: The SQL 1999 Syntax works with other RDBMS which support ANSI SQL

#### SUMMARY:

A join between two tables that return rows that match the join condition and also unmatched rows from left table is LEFT OUTER JOIN

A join between two tables that return rows that match the join condition and unmatched rows from the right table is RIGHT OUTER JOIN

A join between two tables that return rows that match the join condition and returns unmatched rows of both left and right table is a FULL OUTER JOIN

=====  
SubQueries  
=====

We have the concept of INNER Query and OUTER Query.

The INNER Query (sub-query) is executed and the result of it used by the

OUTER Query (main-query)

Getting a list of employees who's salary is more than that of Abel's salary

```
SQL> SELECT first_name FROM employees
      2 WHERE salary > ( SELECT salary FROM employees WHERE last_name =
      'Abel');
```

FIRST\_NAME  
-----

Steven  
Neena  
Lex  
Nancy  
John  
Karen  
Alberto  
Lisa  
Michael  
Shelley

The INNER Query will find Abel's salary and that salary is used as input to the OUTER Query.

Example-2

```
SQL> -- Getting a list of employees who are in 'Finance' department
```

```
SQL>
```

```
SQL> SELECT first_name FROM employees
```

```
2 WHERE department_id = ( SELECT department_id FROM departments WHERE  
department_name = 'Finance' );
```

```
FIRST_NAME
```

```
-----
```

```
Nancy
```

```
Daniel
```

```
John
```

```
Ismael
```

```
Jose Manuel
```

```
Luis
```

Multi-Level Nesting of Subqueries

=====

We want to find the full name of the employee who has joined the company first.

Example-1 : Doing it step by step

```
-- #1 : Getting the lowest HIRE_DATE
```

```
SQL> SELECT MIN(hire_date) FROM employees;
```

```
MIN(HIRE_
```

```
-----
```

```
13-JAN-01
```

```
-- #2 : Getting employee ID of the first employee to hire
```

```
SQL> SELECT employee_id FROM employees WHERE hire_date = '13-JAN-01';
```

```
EMPLOYEE_ID
```

```
-----
```

```
102
```

```
-- #3 : Getting the full name of the employee based on the above ID
```

```
SQL> SELECT first_name || ' ' || last_name "Full Name"
```

```
2 FROM employees
```

```
3 WHERE employee_id = 102;
```

```
Full Name
```

```
-----
```

```
Lex De Haan
```

The same can be achieved by multi-level nesting of queries as follows:

```
SELECT first_name || ' ' || last_name "Full Name"
```

```
FROM employees
```

```
WHERE employee_id = ( SELECT employee_id FROM employees
```

```
WHERE hire_date = (SELECT MIN(hire_date) FROM  
employees ));
```

What happens when a subquery returns multiple rows?

~~~~~

We get an ERROR as follows:

Example:

```
SELECT last_name, salary FROM employees
WHERE salary > ( SELECT salary FROM employees WHERE last_name = 'Smith'
);
```

```
WHERE salary > ( SELECT salary FROM employees WHERE last_name = 'Smith' )
                *
```

ERROR at line 2:

ORA-01427: single-row subquery returns more than one row

To overcome the above error, we need to use MULTI-ROW operators.

Using Multi-Row Operator with SubQueries

-----

Example

SQL> ed

Wrote file afiedt.buf

```
1  SELECT first_name FROM employees
2* WHERE salary IN ( SELECT salary FROM employees WHERE first_name =
'Peter')
SQL> /
```

FIRST\_NAME

-----

Randall  
Martha  
Peter  
Joshua  
James  
Karen  
Hermann  
Harrison  
Janette  
Peter  
Allan  
Peter  
Daniel  
Alexander

NOTE: As we have THREE employees with the first name as 'Peter' we get THREE different salary values.

Thus to get a list of other employees who get a salary similar to that of 'Peter' we use MULTI-ROW operator 'IN' in this case.

Example

```
SQL> ed
Wrote file afiedt.buf
```

```
  1  SELECT first_name FROM employees
  2* WHERE salary NOT IN ( SELECT salary FROM employees WHERE first_name
= 'Peter')
SQL> /
```

To get the list of employees who's salary is not like 'Peter' we can use NOT IN multi-row operator.

NOTE: Topics on the following are coming up in Advance SQL

- Co-related Subqueries
- EXISTS and NOT EXISTS Operators

Set Operators - TDB