

Lesson-04 : Constraints, Adv. Subqueries & Other Database Objects

Constraints

Constraints enforce Business Rules
Constraints prevent deletion of a table if there are dependencies

Constraint Types

PRIMARY KEY
NOT NULL
UNQIUE
CHECK
FOREIGN KEY -> AKA Referential Integrity Constraint

Defining Constraints

Can be defined at COLUMN Level or TABLE Level

[1] COLUMN Level - While creating the DB table.

```
CREATE TABLE <table_name>
(
    colname1    datatype1    Column_CONSTRAINT,
    :
);
```

OR

[2] TABLE Level - While creating the DB table.

```
CREATE TABLE <table_name>
(
    colname1    ..... ,
    colname2    ..... ,
    :
    :
    :
    CONSTRAINT emp_pk PRIMARY KEY( empID )
);
```

If we have forgotten to specify the CONSTRAINT while creating the table we can include constraint with the ALTER TABLE command.

PRIMARY KEY Constraint

Ensures that each and every row/record is identified uniquely.
NOT NULL is IMPLICITLY included in this constraint.

```
ALTER TABLE table2
ADD CONSTRAINT coll_pk PRIMARY KEY (coll);
```

NOT NULL Constraint

~~~~~

Ensure THERE IS data in the column.  
i.e. the column data CANNOT BE EMPTY.

If we are using ALTER TABLE command to specify the NOT NULL constraint, we need to MODIFY the column as the ADD CONSTRAINT clause will not work.

Example:

```
ALTER TABLE customer
MODIFY custCity VARCHAR2(10) CONSTRAINT city_nn NOT NULL;
```

#### UNIQUE Constraint

~~~~~

Ensure the column has UNIQUE values.
However enables to have NULL value(s).

Example:

```
ALTER TABLE customer
ADD CONSTRAINT email_uk UNIQUE ( emailID );
```

CHECK Constraint

~~~~~

Helps in validating the data against a condition.

Example:

```
ALTER TABLE customer
ADD CONSTRAINT age_ck CHECK ( custAge >= 18 );

ALTER TABLE customer
ADD CONSTRAINT gender_ck CHECK (custGender = 'M' or custGender = 'F')
```

#### FOREIGN KEY Constraint

~~~~~

Helps in REFERENTIAL Integrity.

i.e. A foreign key column in one table refers to the primary key column in another table.

Thus any violation will not be accepted.

```
ALTER TABLE table2
ADD CONSTRAINT ms_fk FOREIGN KEY (mStatus) REFERENCES table3 (mStatus);
```

Viewing Constraints

=====

To view the constraints of a table, use the 'user_constraints' data dictionary.

Example:

```
SQL> desc user_constraints
```

```
SQL> SELECT constraint_name, constraint_type, search_condition
2 FROM user_constraints WHERE table_name = 'CUSTOMER';
```

CONSTRAINT_NAME	CONSTRAINT	SEARCH_CONDITION
SYS_C0011124	C	"CUSTNAME" IS NOT NULL
SYS_C0011125	P	
CITY_NN	C	"CUSTCITY" IS NOT NULL
EMAIL_UK	U	
AGE_CK	C	custage >= 18
GENDER_CK	C	custgender = 'M' OR custgender = 'F'

NOTE: In CONSTRAINT_NAME column, any constraint name that starts with SYS_Cxxxxxx is the name given by Oracle Database Server.

However, if you give it explicitly it is more clear and better.

Follow coding style, PRIMARY KEY constraint ends with PK,
 NOT NULL constraint name ends with NN
 UNIQUE constraint name ends with UK
 CHECK constraint name ends with CK
 FOREIGN KEY constraint name ends with FK

Disable Constraint

=====

Execute the DISABLE clause of the ALTER TABLE statement to deactivate an integrity constraint.

Example:

```
SQL> ALTER TABLE customer
  2  DISABLE CONSTRAINT email_uk;
```

```
SQL> INSERT INTO customer VALUES (4, 'Deepak', 'M', 'smartguy@gmail.com',
28, 'Pilani');
```

Enabling Constraint

=====

Activate an integrity constraint that is currently disabled.
 Using the ENABLE clause of the ALTER TABLE statement.

Example:

```
SQL> ALTER TABLE customer
  2  ENABLE CONSTRAINT email_uk;
```

=====

Adv. SubQueries

=====

Multi-Column Subqueries

=====

In multi-column subqueries, each row of the main query is compared to values from multiple-row and multiple-column subquery.

Example:

```
WHERE (manager_id, department_id) IN
```

```

Subquery
    100    90
    102    60
    124    50

```

Column Comparisons

~~~~~

Column Comparisons in a multi-column subquery can be

- [1] Pairwise comparisons
- [2] Non-pairwise comparisons

- [1] Pairwise Comparisons
- ~~~~~

Example:

Displaying the details of employees who are managed by the same manager and work in the same department as the employees with employee ID 174 and 178

```

SQL> SELECT first_name, manager_id, department_id
2   FROM employees
3   WHERE (manager_id, department_id) IN
4         ( SELECT manager_id, department_id FROM employees
5           WHERE employee_id IN (174, 178) )
6   AND employee_id NOT IN (174, 178);

```

| FIRST_NAME | MANAGER_ID | DEPARTMENT_ID |
|------------|------------|---------------|
| Alyssa     | 149        | 80            |
| Jonathon   | 149        | 80            |
| Jack       | 149        | 80            |
| Charles    | 149        | 80            |

- [2] Non-Pairwise Comparisons
- ~~~~~

Example:

Display the details of employees who are managed by the same manager as the employees with ID 174 OR 141 and work in the same department as the employees with ID 174 or 141

Example:

```

SELECT first_name, manager_id, department_id
FROM employees
WHERE manager_id IN
      ( SELECT manager_id FROM employees WHERE employee_id IN (174,141)
      )
AND department_id IN
      ( SELECT department_id FROM employees WHERE employee_id IN
(174,141) )
AND employee_id NOT IN (174, 141);

```

| FIRST_NAME | MANAGER_ID | DEPARTMENT_ID |
|------------|------------|---------------|
| Douglas    | 124        | 50            |
| Donald     | 124        | 50            |
| Kevin      | 124        | 50            |
| Alana      | 124        | 50            |

|          |     |    |
|----------|-----|----|
| Peter    | 124 | 50 |
| Randall  | 124 | 50 |
| Curtis   | 124 | 50 |
| Charles  | 149 | 80 |
| Jack     | 149 | 80 |
| Jonathon | 149 | 80 |
| Alyssa   | 149 | 80 |

#### Scalar Subquery Expression

~~~~~

SCALAR = Single Value

A Scalar Subquery Expression is a subquery that returns exactly one column value for one row.

Can be used in:

- DECODE and CASE condition expressions
- All clauses of SELECT except GROUP BY clause
- The SET clause and the WHERE clause of an UPDATE statement

Example - Using a Subquery in the CASE function

```
SELECT employee_id, last_name,
( CASE
    WHEN department_id = (SELECT department_id FROM departments
                        WHERE location_id = 1800) THEN 'Canada'
    ELSE 'USA'
  END
) Location
FROM employees;
```

EMPLOYEE_ID	LAST_NAME	LOCATION
100	King	USA
101	Kochhar	USA
102	De Haan	USA
103	Hunold	USA
104	Ernst	USA
201	Hartstein	Canada
202	Fay	Canada
203	Mavris	USA
204	Baer	USA

Correlated Subqueries

=====

Correlated subqueries are used for row-by-row processing.
Each subquery is executed once for every row in the outer/main query.

The steps in which the processing of a Correlated Subquery works is as follows:

- [1] GET - Candidate row from the outer/main query
|
- [2] EXECUTE - inner query (subquery) using candidate row value
|
- [3] USE - values from the inner query to qualify or disqualify candidate row.

Syntax:

```
SELECT column1, column2, ...
FROM table1 outer
WHERE column1 operator
      (SELECT column1, column2
       FROM table2
       WHERE expr1 = outer.expr2);
```

The subquery REFERENCE a column from a table in the parent/main query.

Example:

Find all employees who earn more than the average salary in their department.

```
SELECT last_name, salary, department_id
FROM employees outer
WHERE salary > ( SELECT AVG(salary)
                 FROM employees
                 WHERE department_id = outer.department_id);
```

```
SQL> SELECT last_name, salary, department_id
      2 FROM employees outer
      3 WHERE salary > ( SELECT AVG(salary)
      4 FROM employees
      5 WHERE department_id = outer.department_id);
```

LAST_NAME	SALARY	DEPARTMENT_ID
King	24000	90
Hunold	9000	60
Ernst	6000	60
Greenberg	12008	100
Faviet	9000	100
Raphaely	11000	30

Using EXISTS Operator

=====

The EXISTS operator tests for existence of rows in the results set of the subquery.

If a subquery row value is found:

- The search does not continue in the inner query
- The condition is flagged TRUE

If a subquery row value is not found:

- The condition is flagged FALSE
- The search continues in the inner query

Example:

Find employees who have at least one person reporting to them.

```
SELECT employee_id, last_name, job_id, department_id
FROM employees outer
WHERE EXISTS (SELECT 'X' FROM employees WHERE manager_id =
outer.employee_id);
```

```
SQL> SELECT employee_id, last_name, job_id, department_id
2 FROM employees outer
3 WHERE EXISTS (SELECT 'X' FROM employees WHERE manager_id =
outer.employee_id);
```

EMPLOYEE_ID	LAST_NAME	JOB_ID	DEPARTMENT_ID
100	King	AD_PRES	90
101	Kochhar	AD_VP	90
102	De Haan	AD_VP	90
103	Hunold	IT_PROG	60
108	Greenberg	FI_MGR	100
114	Raphaely	PU_MAN	30
120	Weiss	ST_MAN	50

Using NOT EXISTS Operator

=====

Find all departments that do not have any employee.

Example:

```
SELECT department_id, department_name
FROM departments d
WHERE NOT EXISTS (SELECT 'X' FROM employees WHERE department_id =
d.department_id);
```

```
SQL> SELECT department_id, department_name
2 FROM departments d
3 WHERE NOT EXISTS (SELECT 'X' FROM employees WHERE department_id =
d.department_id);
```

DEPARTMENT_ID	DEPARTMENT_NAME
120	Treasury
130	Corporate Tax
140	Control And Credit
150	Shareholder Services
160	Benefits
170	Manufacturing

Using Correlated UPDATE

=====

The syntax for correlated update is as follows:

Syntax:

```
UPDATE table1 alias1
SET column = ( SELECT expression
FROM table2 alias2
WHERE alias1.column = alias2.column );
```

If the subquery reference a column of the parent query for updation then we term this as Correlated Update.

Example

- Denormalize the EMPLOYEES table by adding a column to store the department name.
- Populate the table by using a correlated update.

```
ALTER TABLE emp_temp
ADD(department_name VARCHAR2(16));
```

```
UPDATE emp_temp e
SET department_name =
    ( SELECT department_name
      FROM departments d
      WHERE e.department_id = d.department_id );
```

Correlated DELETE
=====

If the subquery reference a column of the parent query for deletion then we term this as Correlated Delete.

The syntax for using the correlated delete is as follows:

```
DELETE FROM table1 alias1
WHERE column operator
    ( SELECT expression
      FROM table2 alias2
      WHERE alias1.column = alias2.column );
```

Use a correlated subquery to delete rows in one table based on rows from another table.

Example:

Use a correlated subquery to delete only those rows from the EMPLOYEES table that also exist in the EMP_HISTORY table.

```
DELETE FROM employees E
WHERE employee_id =
    ( SELECT employee_id
      FROM emp_history
      WHERE employee_id = E.employee_id );
```

=====

Database Objects

=====

Object	Description
Table	Basic unit of storage, composed of rows & columns
View	Logically represents sub set of data from one or more tables
Sequence	Generate a sequence for primary key values
Index	Improve the performance of queries
Synonym	Alternative name for an object

What is a View?

~~~~~

Refer to PPTs

- NOTE: It is referred a Virtual Table



Why use Views?

~~~~~

- To restrict data access
- To make complex queries easy
- To provide data independence
- To present different views of the same data

Simple v/s Complex Views

~~~~~

| Features               | Simple | Complex     |
|------------------------|--------|-------------|
| No. of table           | One    | One or more |
| Contains functions     | No     | Yes         |
| Contains group of data | No     | Yes         |
| DML operations         | Yes    | No          |

Creating a View

~~~~~

Done using CREATE VIEW statement.

We embed a subquery within a CREATE VIEW statement.

The subquery can contain complex SELECT syntax

Example-1:

```
SQL> CREATE VIEW empvu1
  2 AS SELECT employee_id, first_name, last_name, salary
  3 FROM employees WHERE department_id IN (100, 110);
```

Example-2:

```
SQL> CREATE VIEW salvu50
  2 AS SELECT employee_id ID_NUMBER, last_name NAME, salary*12
ANN_SALARY FROM employees
  3 WHERE department_id = 50;
```

View created.

```
SQL> select * from salvu50;
```

ID_NUMBER	NAME	ANN_SALARY
120	Weiss	96000
121	Frapp	98400
122	Kaufling	94800
123	Vollman	78000
124	Mourgos	69600
125	Nayer	38400

Example-3 : Complex View

```
CREATE VIEW deptsumsal_vu ( dname, minsal, maxsal, avgsal )
AS SELECT d.department_name, MIN(e.salary), MAX(e.salary), AVG(e.salary)
FROM employees e, departments d
WHERE e.department_id = d.department_id
GROUP BY d.department_name;
```

View created.

```
SQL> SELECT * FROM deptsumsal_vu;
```

DNAME	MINSAL	MAXSAL	AVGSAL
Administration	4400	4400	4400
Accounting	8300	12008	10154
Purchasing	2500	11000	4150
Human Resources	6500	6500	6500
IT	4200	9000	5760
Public Relations	10000	10000	10000
Executive	17000	24000	19333.3333
Shipping	2100	8200	3475.55556
Sales	6100	14000	8955.88235
Finance	6900	12008	8601.33333
Marketing	6000	13000	9500

NOTE: The above created view is a complex view, inasmuch it uses the GROUP functions and the GROUP BY clause; besides accessing the data from two tables.

Rules for Performing DML Operations on Views

~~~~~

- You can perform DML operations on simple views.
- You cannot remove a row if the view contains the following:
  - Group functions
  - A GROUP BY clause
  - The DISTINCT keyword
  - The pseudo column ROWNUM keyword
- You cannot modify data in a view if it contains:
  - Group functions
  - A GROUP BY clause
  - The DISTINCT keyword
  - The pseudo column ROWNUM keyword
  - Columns defined by expressions
- You cannot add data through a view if the view includes:
  - Group functions
  - A GROUP BY clause
  - The DISTINCT keyword
  - The pseudo column ROWNUM keyword
  - Columns defined by expressions
  - NOT NULL columns in the base tables that are not selected by the view

#### Inline Views

=====

- An inline view is a subquery with an alias (or correlation name) that you can use within a SQL statement.
- A named subquery in the FROM clause of the main query is an example of an inline view.
- An inline view is NOT a SCHEMA Object.
- Column of inline view can be used in outer query

Example:

List those employees whose salary is more than the average salary of employees

```
SQL> SELECT first_name || ' ' || last_name "NAME", Salary, AvgSalary
  2   FROM employees, (SELECT AVG(salary) AvgSalary FROM employees)
  3  WHERE salary > avgsalary;
```

| NAME          | SALARY | AVGSALARY  |
|---------------|--------|------------|
| Steven King   | 24000  | 6461.83178 |
| Neena Kochhar | 17000  | 6461.83178 |
| Lex De Haan   | 17000  | 6461.83178 |

Using the WITH CHECK OPTION Clause

=====

- You can ensure that DML operations performed on the view stay within the domain of the view by using the WITH CHECK OPTION clause.
- Any attempt to change the department number for any row in the view fails because it violates the WITH CHECK OPTION constraint.

Example:

```
SQL> CREATE VIEW empvu110
  2   AS SELECT * FROM employees WHERE department_id = 110
  3   WITH CHECK OPTION CONSTRAINT empvu110_ck;
```

If we try to change the department\_id from 110 to any other department\_id we get the following error:

```
SQL> UPDATE empvu110
  2   SET department_id = 100;
UPDATE empvu110
      *
ERROR at line 1:
ORA-01402: view WITH CHECK OPTION where-clause violation
```

Denying DML Operations

=====

- You can ensure that no DML operations occur by adding the WITH READ ONLY option to your view definition.
- Any attempt to perform a DML on any row in the view results in an Oracle server error.

Example:

```
SQL> CREATE VIEW empvu10
  2   AS SELECT * FROM employees WHERE department_id = 10
  3   WITH READ ONLY;
```

```
SQL> DELETE FROM empvu10;
DELETE FROM empvu10
      *
ERROR at line 1:
ORA-42399: cannot perform a DML operation on a read-only view
```

## Removing a View

=====

To remove a view use the DROP VIEW statement.

Example:

```
SQL> DROP VIEW salvu50;
```

View dropped.

## Indexes

=====

What is an Index?

~~~~~

An index:

- Is a schema object
- Is used by the Oracle server to speed up the retrieval of rows by using a pointer
- Can reduce disk I/O by using a rapid path access method to locate data quickly
- Is independent of the table it indexes
- Is used and maintained automatically by the Oracle server

How are Indexes Created?

~~~~~

- AUTOMATICALLY: A unique index is created automatically when you define a PRIMARY KEY or UNIQUE constraint in a table definition.
- MANUALLY: Users can create non-unique indexes on columns to speed up access to the rows.

Creating an Index

~~~~~

Done using CREATE INDEX statement

Syntax:

```
CREATE INDEX index_name  
ON table (column[, column]...);
```

Example:

```
SQL> CREATE INDEX emp_lastname_idx  
2 ON employees (last_name);
```

Index created.

When to Create an Index?

~~~~~

You should create an index if:

- A column contains a wide range of values
- A column contains a large number of null values
- One or more columns are frequently used together in a WHERE clause or a join condition
- The table is large and most queries are expected to retrieve less than 2 to 4 percent of the rows

When NOT to Create an Index?

~~~~~

It is usually not worth creating an index if:

- The table is small
- The columns are not often used as a condition in the query
- Most queries are expected to retrieve more than 4 percent of the rows
- The table is updated frequently
- The indexed columns are referenced as part of an expression

Removing an Index

~~~~~

To remove an Index, the DROP INDEX statment is used.

```
SQL> DROP INDEX emp_lastname_idx;
```

Index dropped.

=====

Sequences

=====

What is a Sequence?

~~~~~

A sequence:

- Automatically generates unique numbers
- Is a SHARABLE OBJECT
- Is typically used to create a primary key value
- Replaces application code
- Speeds up the efficiency of accessing sequence values when cached

Creating a Sequence

~~~~~

To create a sequence the CREATE SEQUENCE statement is used.

Syntax:

```
CREATE SEQUENCE sequence
[INCREMENT BY n]
[START WITH n]
[{MAXVALUE n | NOMAXVALUE}]
[{MINVALUE n | NOMINVALUE}]
[{CYCLE | NOCYCLE}]
[{CACHE n | NOCACHE}];
```

Example:

```
SQL> CREATE SEQUENCE my_seq
 2  START WITH 5
 3  INCREMENT BY 2
 4  MAXVALUE 99
 5  NOCACHE
 6  NOCYCLE;
```

Sequence created.

Using a Sequence

~~~~~

NEXTVAL and CURRVAL Pseudo columns

- NEXTVAL returns the next available sequence value.
It returns a unique value every time it is referenced, even for different users.

- CURRVAL obtains the current sequence value.
- NEXTVAL must be issued for that sequence before CURRVAL contains a value.

Example: Use sequence 'my_seq' NEXTVAL in the INSERT statement.

```
SQL> INSERT INTO table2 VALUES(my_seq.NEXTVAL, 'Raju', 'M', 'S');
```

1 row created.

```
SQL> INSERT INTO table2 VALUES(my_seq.NEXTVAL, 'Rani', 'F', 'M');
```

1 row created.

```
SQL> SELECT * FROM table2;
```

COL1	COL2	C	M
5	Raju	M	S
7	Rani	F	M

- View the current value for the 'my_seq' sequence.

```
SQL> SELECT my_seq.CURRVAL FROM dual;
```

CURRVAL
7

Since, the last inserted value was 7, we are able to get that value

Using a Sequence

~~~~~

- Caching sequence values in memory gives faster access to those values.
- Gaps in sequence values can occur when:
  - A rollback occurs
  - The system crashes
  - A sequence is used in another table
- If the sequence was created with NOCACHE, view the next available value, by querying the USER\_SEQUENCES table.

Modify a Sequence

~~~~~

Change the increment value, maximum value, minimum value, cycle option, or cache option.

Example:

```
ALTER SEQUENCE my_seq
INCREMENT BY 5
MAXVALUE 999
CACHE
NOCYCLE;
```

Removing a Sequence

~~~~~

- Remove a sequence by using the DROP SEQUENCE statement.
  - Once removed, the sequence can no longer be referenced.
- ```
DROP SEQUENCE my_seq;
```

## Synonyms

Simplify access to objects by creating a synonym (another name for an object).

With synonyms, you can:

- Ease referring to a table owned by another user
- Shorten lengthy object names

Syntax:

```
CREATE [PUBLIC] SYNONYM synonym
FOR object;
```

## Creating and Removing Synonyms

- Create a shortened name for the MY\_SEQ sequence.

Example:

```
SQL> CREATE SYNONYM ms FOR my_seq;
```

Checking the synonym

```
SQL> SELECT ms.CURRVAL FROM dual;
```

```
      CURRVAL
-----
          9
```

- Drop a synonym.

Example:

```
SQL> DROP SYNONYM ms;
```

## Multi-Row Subqueries

- Return more than one row
- Use multiple-row comparison operators

## Multi-Row Operators

| Operator | Description                                           |
|----------|-------------------------------------------------------|
| IN       | Equal to any member in the list                       |
| ANY      | Compare value to each value returned by the subquery  |
| ALL      | Compare value to every value returned by the subquery |

NOTE:

- < ANY Means, LESSER than the HIGHEST value returned by the subquery
- < ALL Means, LESSER than the LOWEST value returned by the subquery

To understand the above, let us first get the salary values for 'MANAGER'

```
SQL> SELECT sal FROM emp WHERE job = 'MANAGER';
```

| SAL  |
|------|
| 2975 |
| 2850 |
| 2450 |

Examples:

```
SQL> SELECT ename, sal FROM emp
2 WHERE sal < ANY ( SELECT sal FROM emp WHERE job = 'MANAGER' );
```

| ENAME  | SAL  |
|--------|------|
| SMITH  | 800  |
| JAMES  | 950  |
| ADAMS  | 1100 |
| WARD   | 1250 |
| MARTIN | 1250 |
| MILLER | 1300 |
| TURNER | 1500 |
| ALLEN  | 1600 |
| CLARK  | 2450 |
| BLAKE  | 2850 |

```
SQL> SELECT ename, sal FROM emp
2 WHERE sal < ALL ( SELECT sal FROM emp WHERE job = 'MANAGER' );
```

| ENAME  | SAL  |
|--------|------|
| ALLEN  | 1600 |
| TURNER | 1500 |
| MILLER | 1300 |
| WARD   | 1250 |
| MARTIN | 1250 |
| ADAMS  | 1100 |
| JAMES  | 950  |
| SMITH  | 800  |